COMPARISON · LAYER 3 · DEPLOY AND SHARE

Cloud Studio vs OpenShift

describe the deployment, it ships.

Red Hat OpenShift (IBM) is a broad platform your team runs. Cloud Studio is an agent that provisions and deploys for you.

Shayak Mazumder
Founder, CEO and CTO, Adya AI

RED HAT OPENSHIFT

A broad Kubernetes platform your team operates.



vs

ADYA · CLOUD STUDIO

An agent that provisions and deploys on three clouds.

AGENTIC MODE

Two ways to ship to the cloud: *run a platform, or brief an agent.*

OPENSHIFT · YOUR TEAM RUNS IT

A large, proven platform. Your platform team operates it.

01 Design clusters, networking and storage

02 Write manifests, pipelines and policies

03 Patch, upgrade and troubleshoot

CLOUD STUDIO · THE AGENT RUNS IT

You describe what to deploy. The agent plans and ships it.

01 Provisions VMs, Kubernetes, serverless and GPUs

02 Sets up CI/CD and deploys

03 Compares live prices across clouds and regions

OpenShift Lightspeed is an AI assistant that answers and troubleshoots. It does not plan or provision for you.

Where OpenShift is bigger, *and where Cloud Studio is different.*

WHERE OPENSHIFT LEADS

A far broader platform
Service mesh, virtualization, GitOps, pipelines and multi-cluster management.

Enterprise depth
Years of support, certifications and a large partner ecosystem.

Runs almost anywhere
Managed on AWS, Azure, IBM Cloud and Google Cloud, or on-prem.

WHERE CLOUD STUDIO IS DIFFERENT

Agentic from the first prompt
It picks instance types, regions and GPUs with you, then provisions them.

Focused on purpose
Provisioning, deployment and CI/CD, done for you rather than run by a team.

Built into the fabric
SAI, Model Studio and every other studio deploy through it.

Cloud Studio and OpenShift, *side by side.*

	Red Hat OpenShift	Adya Cloud Studio
Scope	Full app platform, mesh, virtualization	Provisioning, deployment and CI/CD
Who runs it	A platform team, with operators	An agent from a prompt; you approve
AI help	Lightspeed assistant answers questions	Plans, provisions and deploys itself
Clouds	AWS, Azure, IBM, Google, on-prem	AWS, Azure and Google Cloud, one console
Regions	Per cluster you build	40+ data centres, region choice
Compute	Containers and VMs	VMs, Kubernetes, serverless and GPUs
GPUs	Through OpenShift AI and operators	GPU-as-a-service, single or multi-node
Cost view	Subscription plus cloud bills	Live pricing across clouds and regions

One app, two journeys: *from code to production.*

OPENSHIFT

Your platform team

01	02	03	04	05	06
Size and install clusters	Set up network and storage	Write manifests or Helm	Build pipelines	Deploy and observe	Patch and upgrade

ADYA · CLOUD STUDIO

The agent, you approve

01	02	03	YOU	04	05	06
Reads the app and goal	Compares clouds and prices	Proposes the plan	**You approve**	Provisions in your tenancy	Sets up CI/CD, deploys	Monitors and scales

The agent does about 95% of the work. You approve the plan and the spend.

Three clouds, *one agent, your tenancy.*

3

Clouds, one console

AWS, Azure and Google Cloud

OPENSHIFT
A managed edition per cloud

40+

Data centres

Pick the region; data stays where you put it

OPENSHIFT
Region set per cluster you build

2

Planes

Adya runs the management plane. Your workloads and data stay in your cloud.

OPENSHIFT
You run the whole platform

What Cloud Studio does *on its own.*

WHAT IT DOES

01 AWS, Azure and Google Cloud, one console

02 40+ data centres, with region choice

03 SaaS, PaaS, private cloud, on-prem or edge

04 VMs, Kubernetes, serverless and GPUs

05 CI/CD, observability and logs built in

06 Live pricing across clouds and regions

WHY IT MATTERS

Your data stays in your tenancy

Adya runs only the management plane. Workloads and data live in your account.

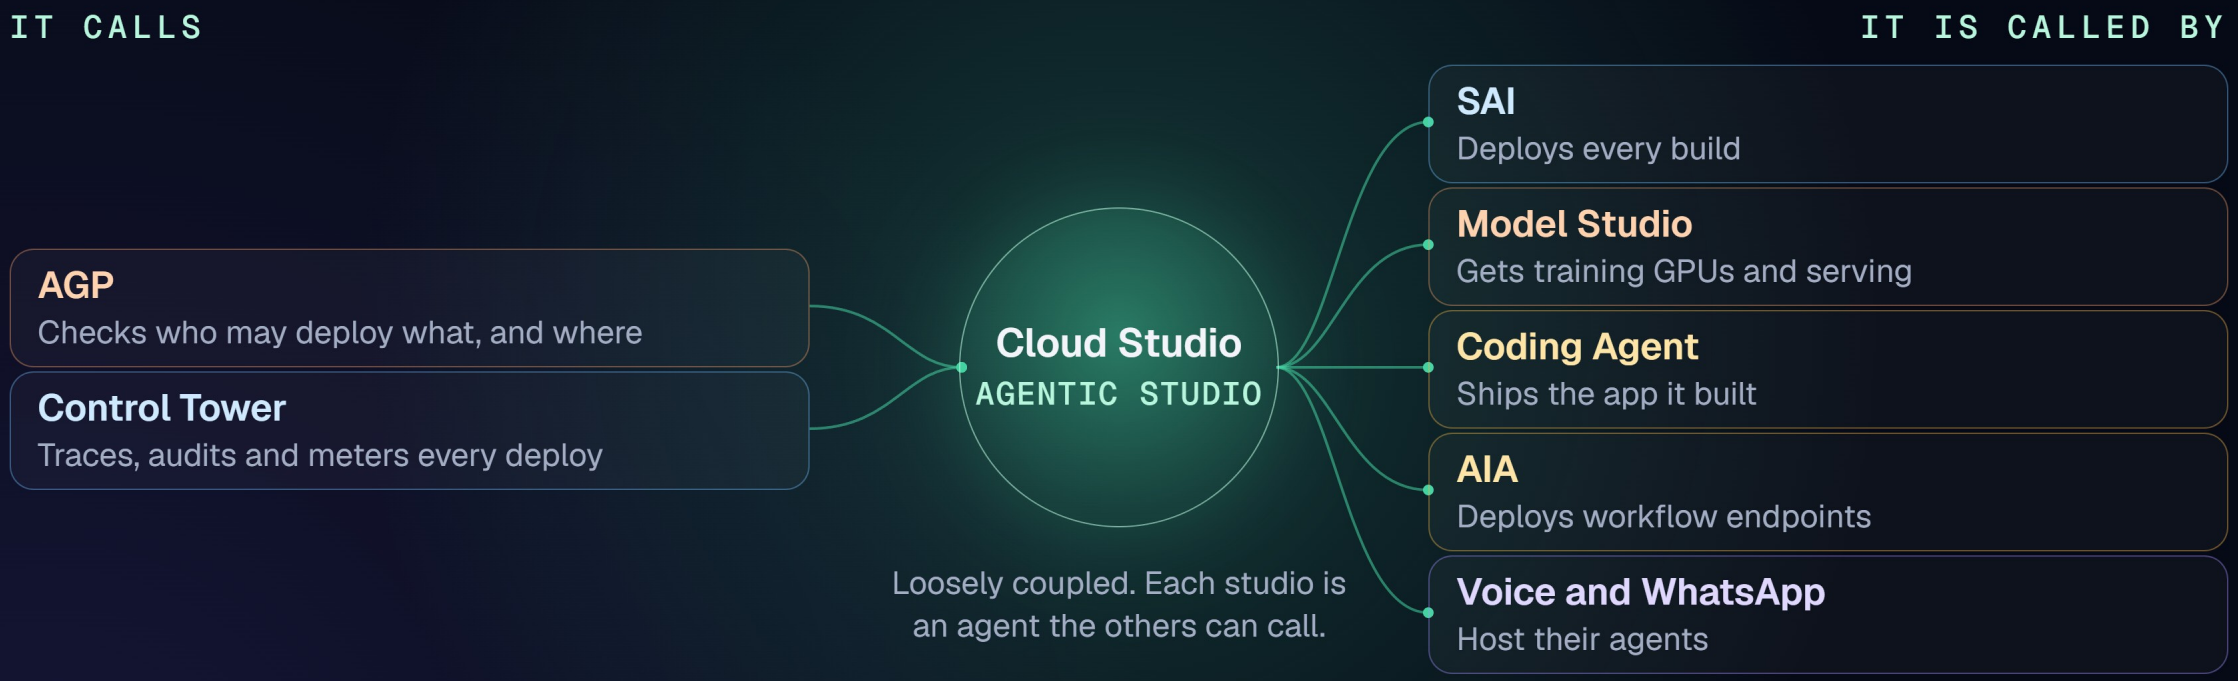
GPUs on demand

Training environments, including multi-node, as GPU-as-a-service, priced across clouds.

Developer mode

Every plan is inspectable and editable. Run any step by hand when you want to.

Every studio deploys *through Cloud Studio.*



Under every studio, *the same three IPs.*

ESM · EVENT-SOURCED MEMORY

Remembers everything

Remembers each deployment, its plan and its outcome, so the next one is faster and cheaper.

AGP · ADAPTIVE GOVERNANCE

Rules agents cannot break

Deployment rules for regions, budgets and who may ship are checked before anything is provisioned.

MAN HARNESS

Hundreds of agents, no bottleneck

Many deployments and environments run in parallel, with no central bottleneck.

Developer mode

Run any step by hand when you want to. The agent does about 95%; you supervise and decide.

Two planes

Management plane by Adya, user plane in your cloud. No customer workload runs on the management plane.

Alone, it ships the app.
Together, it ships the whole system.

"Fine-tune a support model, wrap it in an app, and run it in our AWS account in Mumbai."

01
Model Studio
Fine-tunes the support model

→

02
Coding Agent
Builds the app and its API

→

03
Cloud Studio
Provisions GPUs and deploys in Mumbai

→

04
AGP
Holds the release to your rules

One request. Four studios. Everything lands in your own cloud account.

Stop operating the platform.

Brief the deployment.

Thank you.

Shayak Mazumder / Founder, CEO and CTO, Adya AI

IN ONE LINE

- **OpenShift**
  A broad platform your team runs
- **Cloud Studio**
  An agent that provisions and deploys
- **Together**
  Build, train and ship on one fabric

EXPLORE

vanij.adya.ai