COMPARISON · LAYER 1 · FOUNDATIONS

Model Studio vs Unsloth

a model team, on demand.

Excellent training engines, against an agentic studio that runs all seven steps, from research to deployment.

Shayak Mazumder
Founder, CEO and CTO, Adya AI

UNSLOTH · TOGETHER AI

Fast engines. A data scientist drives every step.



vs

ADYA · MODEL STUDIO

Agents that research, train, test and deploy it.

AGENTIC MODE

Two ways to get your own model: *hire the team, or brief the studio.*

UNSLOTH · TOGETHER AI

Great engines. A data scientist still drives every step.

01 You pick the base model, recipe and settings

02 You find, clean and format the data

03 You run evals, then export and deploy

MODEL STUDIO · THE AGENTS RUN IT

Agents run all seven steps. You approve spend and release.

01 Researches the task, finds or creates the data

02 Chooses the model, GPU and training recipe

03 Trains, evaluates, rolls back or ships

Closest agentic peer: Hugging Face's ML Intern, an open-source CLI agent tied to the Hugging Face ecosystem.

Seven steps to a model, *run by agents.*

01	02	03	04	05	06	07
Understand the task	**Research**	**Data**	**Choose the model**	**GPU and recipe**	**Train and evaluate**	**Accelerate and deploy**
Turns your brief into a spec, success metrics and a budget.	Reads papers and methods that shape the recipe.	Finds, scrapes, cleans or synthesises the data.	Ranks 171 base models by fit, licence and size.	Picks GPUs, region and training parameters.	Trains with checkpoints, runs evals, rolls back if worse.	vLLM or NIM. CUDA on one GPU, Slurm on many.
`Orchestrator`	`Recipe agent`	`Data agent`	`Model agent`	`Compute agent`	`Training, Eval`	`Cloud Studio`

Two human gates: approve the data and budget before training, and the release before deployment.

Everything else runs on its own. In developer mode you can run any step by hand.

What Unsloth and Together do well, *and where Model Studio keeps going.*

WHERE THEY ARE STRONG

Unsloth: fast and lean

About 2x faster training with up to 70% less VRAM, open source.

Unsloth Studio

A local no-code UI since March 2026: 500+ models, data recipes, LoRA to full tuning.

Together AI

200+ hosted open models, managed fine-tuning, dedicated endpoints and GPU clusters.

WHERE THEY STOP, AND MODEL STUDIO KEEPS GOING

You prompt, Model Studio builds the LLM

Model Studio's agents choose the data, model, GPU and recipe, and explain why.

From Data to Evals

Frozen holdouts and quality gates. A worse candidate is rolled back, not shipped.

GPU as a Service and Acceleration

Model Studio deploys with vLLM or NIM on your cloud and manages scaling.

Model Studio, Unsloth and Together, *side by side.*

	Unsloth	Together AI	Adya Model Studio
Who drives	A data scientist	A data scientist	Agents; you approve
Data	Bring it; Data Recipes help	Upload your dataset	Finds, scrapes, synthesises, cleans
Model choice	You pick from 500+	You pick from the catalogue	Agent ranks 171 base models
Compute	Your GPU or a rented one	Together's GPU cloud	Cloud Studio: AWS, Azure, GCP
Evals and rollback	You script them	You run them	Frozen gates, automatic rollback
Deployment	Export GGUF or adapters	Together endpoints	vLLM or NIM; CUDA or Slurm
Where it runs	Your hardware	Together's cloud	Your own cloud account

Unsloth and Together details from their docs and 2026 coverage, October 2026.

Your own model, *at a fraction of the cost.*

171

Base models

SLMs and LLMs, ranked by fit, licence and size

UNSLOTH STUDIO
500+ models, but you choose

1B–100B

Parameters

From the edge to the cloud

TOGETHER AI
Hosted on Together's cloud

~90%

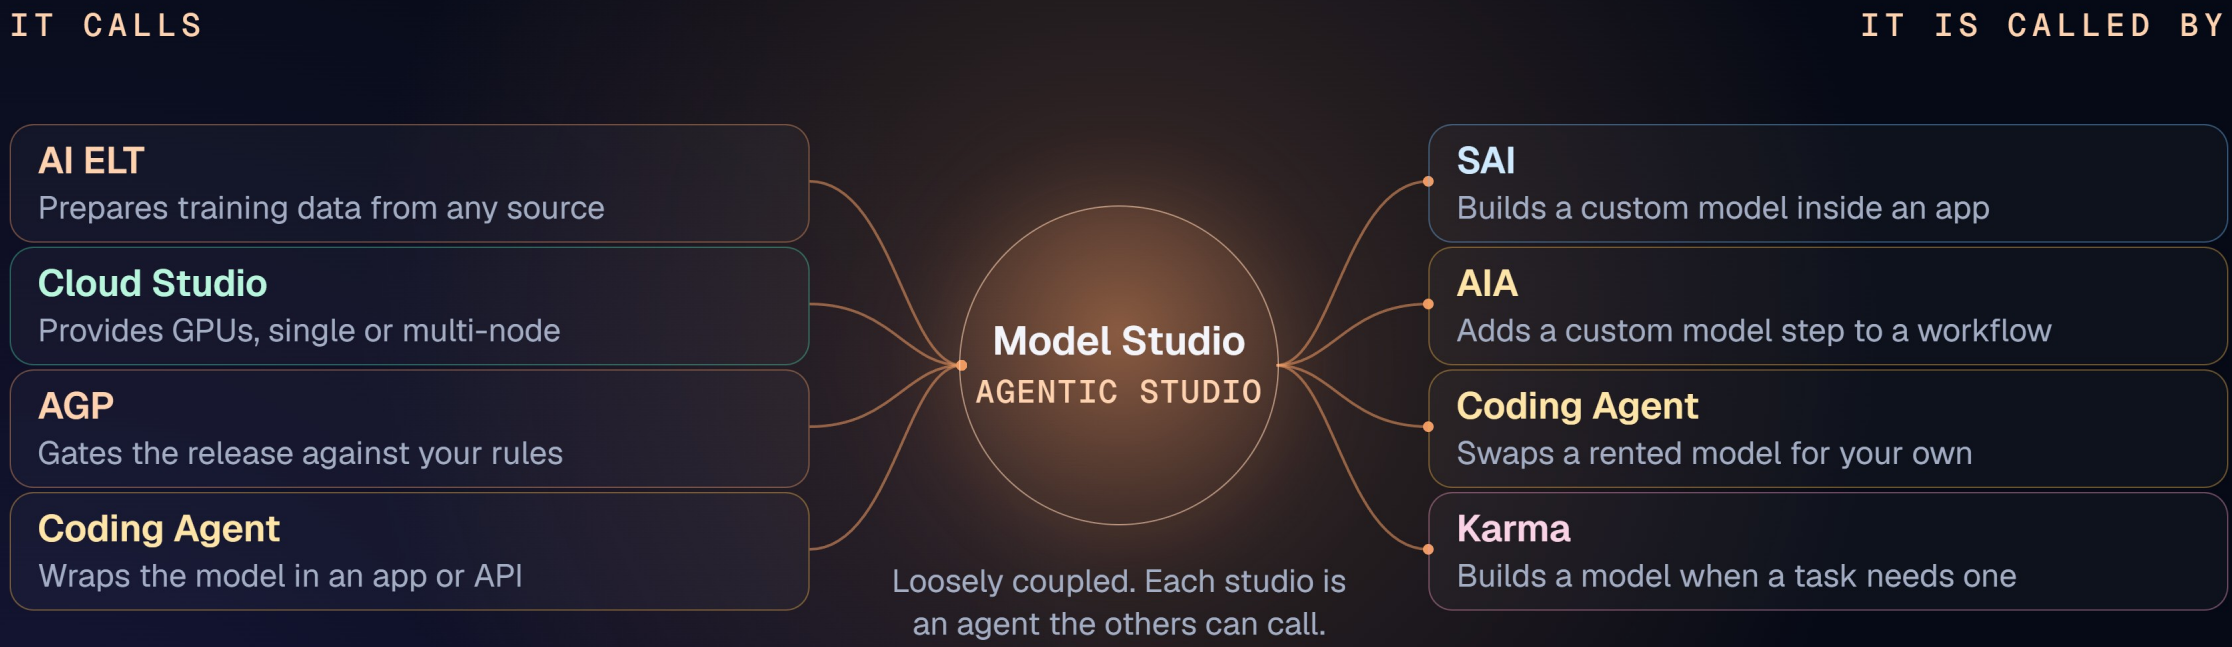
Lower cost

Your own model on solved tasks, against renting a frontier model

RENTED MODELS
Pay per token, forever

ESM trims a further ~70% of tokens across runs, the same layer behind RAS's results against Claude.

Model Studio calls other studios, *and they call it.*



Under every studio, *the same three IPs.*

ESM · EVENT-SOURCED MEMORY

Remembers everything

Remembers every run, recipe and result, so the next model starts from what worked.

AGP · ADAPTIVE GOVERNANCE

Rules agents cannot break

Data permissions, budgets and releases are rules, checked before any GPU time is spent.

MAN HARNESS

Hundreds of agents, no bottleneck

Data, model, compute, training and eval agents work in parallel, peer to peer.



Developer mode

Edit the recipe as YAML and run any step by hand. The agents do about 95%; you supervise.



GPU-as-a-service

Training GPUs through Cloud Studio, from a single machine to multi-node.

Alone, it trains the model.
Together, it ships the product.

"Build a compact contract-review assistant on our clause data, and keep the pilot under $100."

01
AI ELT
Pulls and structures the clause data

→

02
Model Studio
Fine-tunes a compact model with QLoRA

→

03
AGP
Gates the release on accuracy and policy

→

04
Coding Agent
Builds the review app around it

→

05
Cloud Studio
Serves it in your tenancy

One brief. Five studios. A model you own, inside a product you ship.

Stop hiring a model team.

Brief the studio.

Thank you.

Shayak Mazumder / Founder, CEO and CTO, Adya AI

IN ONE LINE

- **Unsloth and Together**
  Engines a data scientist drives
- **Model Studio**
  Agents that run all seven steps
- **Together**
  Data, model, app and cloud on one fabric

EXPLORE

vanij.adya.ai