COMPARISON · LAYER 4 · CHANNELS

Voice Studio vs Vapi

from one prompt to a running campaign.

A voice runtime you assemble and tune, against an agentic studio that plans, builds, tests, runs and learns.

Shayak Mazumder
Founder, CEO and CTO, Adya AI

VAPI

A strong voice runtime you configure and tune.



vs

ADYA · VOICE STUDIO

Nine agents that build, launch and improve it.

AGENTIC MODE

Two ways to build a voice agent: *configure it, or brief it.*

VAPI · YOU ASSEMBLE IT

You pick the stack, write the prompt and wire the rest.

01 Choose transcriber, model and voice per assistant

02 Composer drafts an assistant from a prompt

03 Workflows, integrations and deployment stay with you

VOICE STUDIO · THE AGENTS BUILD IT

You give the brief. Nine agents take it to a tested, live campaign.

01 Chooses ASR, LLM, TTS or speech-to-speech

02 Builds the workflow, tools, knowledge and handoffs

03 Simulates, deploys, analyses and learns

Vapi's Composer drafts an assistant from a prompt. It does not build the full workflow, integrations or deployment on its own.

What Vapi does well, *and where Voice Studio keeps going.*

WHERE VAPI IS STRONG

A mature voice runtime
Low-latency calling with a wide choice of transcribers, models and voices.

Squads for handoffs
Multiple assistants with defined handoffs, run at large call volumes.

Developer-first
Clean APIs and SDKs, MCP tool calling and transparent per-minute pricing.

WHERE IT STOPS, AND VOICE STUDIO KEEPS GOING

Voice AI Agent Builder
The Voice AI agent picks ASR, LLM, TTS or speech-to-speech by language, latency and cost.

Voice Agent Runs Autonomously
VOX agents build the workflow, tools, telephony, tests and campaign end to end.

More than just a Voice Agent
Voice Studio calls AIA, WhatsApp, Coding Agent and Cloud Studio for everything around the call.

Voice Studio and Vapi, *side by side.*

	Vapi	Adya Voice Studio
Build mode	Dashboard, API and Composer	Agentic mode, plus full manual studios
Voice stack	You choose per assistant	Voice AI agent picks by language, latency, cost
Speech	STT, LLM and TTS, or realtime	ASR, LLM and TTS, or speech-to-speech
Indian languages	Through the providers you pick	Indic ASR and TTS, caller-led switching
Workflows	Squads; workflow builder deprecated	Workflow agent builds it on a shared canvas
Testing	Evals; regression and load on you	QA agent tests each exact revision
Campaigns	Outbound via API and campaigns	Consent, budget, call windows and retries built in
Memory	Per-call context	ESM remembers every caller, safely

Vapi details from Vapi's docs, pricing page and 2026 reviews, October 2026.

One brief, two journeys: *a Hindi and English sales line.*

VAPI

You, with Composer's help

01	02	03	04	05	06
Pick transcriber, model, voice	Write prompt and tools	Buy or import numbers	Set up squads and handoffs	Test calls, tune latency	Wire analytics and retries

ADYA · VOICE STUDIO

Nine agents, you approve

01	02	03	04	05	YOU	06
Brief and language	Voice stack	Knowledge and memory	Conversation and workflow	Calls and tests	**You approve**	Runs, analyses, learns

The agents do about 95% of the work. You approve the release and the campaign scope.

Built for volume, *in every language you sell in.*

9

Specialist agents

From brief to QA, plus the VOX orchestrator

VAPI
One Composer copilot

100s

Calls in parallel

Within the budget and concurrency you approve

VAPI
10 lines included, then $10 per line a month

<2s

Speech to speech

Real-time voice, end to end

VAPI
Depends on the stack you assemble

~70%

Fewer tokens

ESM carries context across calls, so turns get shorter

VAPI
Cross-call memory is yours to build

ESM is the same memory layer behind RAS's 3.9x token cut against Claude Life Sciences, published October 2026.

What the VOX agents *decide for you.*

WHAT THEY DECIDE

01 Language policy and code-switching

02 ASR, LLM, TTS or realtime voice per call

03 Knowledge with provenance and scoped memory

04 Prompt, greeting, fillers and turn-taking

05 Tools, webhooks and human handoff on a canvas

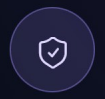
06 Telephony route, call windows and retries

WHY IT MATTERS



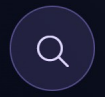
Two modes, one agent

Agentic mode and the manual studios share every version, test and log.



Production controls built in

Noise handling, DTMF, redaction, consent checks and human in the loop.



Analytics that answer

A voice analyst agent reads your calls and explains latency, outcomes and cost.

Voice Studio calls other studios, *and they call it.*

IT CALLS

AIA
Builds the after-call workflow

Coding Agent
Builds the app around the call

WhatsApp Studio
Sends the follow-up journey

RAS
Brings cited research into knowledge

Cloud Studio
Deploys the agent and its stack



Loosely coupled. Each studio is an agent the others can call.

IT IS CALLED BY

RAI
Runs sales and qualification calls

Karma
Makes calls on your behalf

SAI
Puts a voice agent inside any app

AIA
Adds a call step to a workflow

Under every studio, *the same three IPs.*

ESM · EVENT-SOURCED MEMORY

Remembers everything

Remembers each caller and each call. Fewer turns, fewer tokens and a lower cost per minute of talk.

AGP · ADAPTIVE GOVERNANCE

Rules agents cannot break

Controls what the agent may say and do on every call. If a rule fails, the action does not run.

MAN HARNESS

Hundreds of agents, no bottleneck

Hundreds of calls and agents run in parallel, with no central bottleneck.

Developer mode

Run any step by hand when you want to. The agent does about 95%; you supervise and decide.

Human Assistant

Answers routine questions for you from approved sources. Budgets and commitments wait for you.

Alone, Voice Studio takes the call.
Together, it runs the front desk.

"Answer our clinic line in Hindi and English, book visits and confirm them on WhatsApp."

01
Voice Studio
Answers calls in Hindi and English

→

02
AIA
Books the visit and updates the CRM

→

03
WhatsApp Studio
Sends the confirmation and intake form

→

04
Coding Agent
Builds the front-desk dashboard

→

05
Cloud Studio
Deploys it in your tenancy

One brief. Five studios. ESM remembers every patient across voice and chat.

Stop assembling. Start briefing.

Voice Studio runs it.

Thank you.

Shayak Mazumder / Founder, CEO and CTO, Adya AI

IN ONE LINE

- **Vapi**
  A runtime you assemble and tune
- **Voice Studio**
  Agents that build, run and improve it
- **Together**
  Voice, chat, workflow and app on one fabric

EXPLORE

vanij.adya.ai